

CS307 Spring 2021 Database Project 2

1. Source code

Download link:

<https://mirrors.sustech.edu.cn/git/sustech-2021fall-db/sustech-sql-project>

Interface Specification

The structure of the interfaces is as follows.

- `database` folder stores connection information such as username, password, url, we only provides `PostgreSQL` as the DBMS.
- `dto` folder stores a set of data objects that will be accessed by interfaces. **Your implementation will use them as parameters or returned values.**
- `service` folder stores **Service Interfaces**, this is the folder you should pay special attention to. There exist multiple `.java` file where the interface signatures are stored. You need to implement you own `class` to fit these signatures.
- `exception` folder stores exceptions that you should `throw` if something went wrong.
- `factory` folder stores the `ServiceFactory` abstract class that you need to implement to create your service instances.

Your Tasks

- Implement the `service` and `factory` interfaces to pass the base testcases.
- Design your (PostgreSQL) database to satisfy the requirements of interfaces.
- Profile your implementation and find ways to speed it up.
- (Optional) Find other ways to implement similar functionalities as our interfaces and compare (some of) them, are they better, worse or have different use cases.

Here is a reference implementation, it shows you how to implement one method of an interface. To get a service working, you'll have to implement **all its interfaces**

The following code is just a guide, the code interacts with database will usually be written in the DAO layer

```
@ParametersAreNonnullByDefault
public class ReferenceStudentService implements StudentService {
    /* Some codes are omitted */
    @Override
    public void dropCourse(int studentId, int sectionId) {
        try (Connection connection = SQLDataSource.getInstance().getSQLConnection();
            PreparedStatement stmt = connection.prepareStatement("call drop_course(?,
?)")) {
            stmt.setInt(1, studentId);
            stmt.setInt(2, sectionId);
            stmt.execute();
        } catch (SQLException e) {
```

```

        e.printStackTrace();
    }
}
/* Some codes are omitted */
}

```

```

public class ReferenceServiceFactory extends ServiceFactory {
    public ReferenceServiceFactory() {
        registerService(StudentService.class, new ReferenceStudentService());
        registerService(CourseService.class, new ReferenceCourseService());
        // registerService(<interface name>.class, new <your implementation>());
    }
}

```

After you have implemented your factory class, be sure to put your factory class name into the file `./config.properties`. So that we can find your implementation and test.

```

serviceFactory=your.package.YourServiceFactory           // Your factory class name
here.
jdbcUrl=jdbc:postgresql://localhost:5432/project2
username=postgres
password=postgres

```

Tips

Please create database with `c` locale, which provides the platform-independent sorting result.

Here is a sample commands:

```
CREATE DATABASE project2 WITH ENCODING='UTF8' LC_COLLATE = 'C';
```

See <https://github.com/NewbieOrange/SUSTech-SQL-Project2-Public/issues/88>,

and <https://stackoverflow.com/questions/43890221/column-sorting-in-postgresql-is-different-between-macos-and-ubuntu-using-same-co>

Additional requirements of interface

Java

大概

- All `add*()` functions with `int` as return value should return the (presumably auto-generated) ID.
- All arguments are guaranteed to be non-null, unless marked as `@Nullable`.
- All return values (and their fields) should be non-null, unless explicitly documented otherwise. If a `list/map` is empty, put `List.of()/Map.of()` or equivalents instead of `null`.
- Do **NOT** modify anything in the provided interfaces, or any of the framework code.
- Your implementation should throw `java.lang.UnsupportedOperationException` if a method is not actually implemented.

优化查询速度：
在开头加explain & explain analyse

so the tests can fail quickly.

Rules

- Data should be persisted on disk after each write operation instead of only modified in RAM. If you introduced a cache layer, you have to enforce the consistency. You should also ensure the durability in case of a sudden shutdown. 稳定性
- You should **NOT** use frameworks such as **ORM**.
- You don't need to spend time on **GUI/WEB**, as we do **NOT** give extra scores for them.

Java-specific rules

- You should **NOT** modify or add any class in package `cn.edu.sustech.cs307`. Use another package for your implementations.
- You should **NOT** extend any class in package `cn.edu.sustech.cs307.dto`.
- In this project, we use Maven to manage dependent libraries. If you want to introduce a new library, you need to record it in `pom.xml`. Your dependencies should be downloadable from the Maven Central repository.

2. What to deliver?

- **PASS BASE TEST:** First and foremost, you should pass the base testcases, this is the basic requirement.
- **IMPROVE YOUR EFFICIENCY:** After you passed the base tests, you need to find ways to improve the performance of your implementation. You can work on the following aspects.

Resource Consumption

- Memory Consumption: How much memory your database takes?
- Disk Consumption: How much disk space your database takes? How are they distributed? (index, data, other info?)

Speed

- Data Import Speed: How much time your database need to import all data?
- Data Modify Speed (Insertion, Update, Deletion): How much time your database need to change one/one hundred/one million rows?
- Data Query Speed: How much time your database need to fetch one/one hundred/one million rows?
- Cache Policy: How much time your database need to fetch a row if the row was just accessed by others?

Concurrency

- Simultaneous Query Number: How many queries can your database handles simultaneously?
- Simultaneous Query Latency: How long does it take to query if there are many users connect to your database simultaneously.
- Transaction Safety: Is your database safe with many users concurrently writing/reading to it?

Correctness

畸形的

- Malformed Data Identification: Can your database identify the malformed data automatically?
- ACID Principle
- **(Optional) DIFFERENT WAYS SAME GOAL?** Can you find other ways to implement these functionalities? Are they **BETTER/WORSE/USECASE-RELATED?** Please do share us your amazing ideas.

Project Timeline

Code Submission Deadline: **December 31st, 2021 18:30**

Presentation Time: **December 31st, 2021** in Lab class